

**AK**

1. For the following, circle **TRUE** or **FALSE** (24 points)

- a. Two's complement allows us to represent the same ~~number~~ of positive integers as negative integers for a given word length.    **TRUE**    **FALSE**
- b. If `%edx = 0xABABABAB` and `%eax = 0x4455AA22`, the x86-64 instruction  
    `movsbl %dh, %eax`  
    would change the value of `%rax` to `0xFFFFFFFFAB`    **TRUE**    **FALSE**
- c. The variables local to a function are allocated from the stack.    **TRUE**    **FALSE**
- d. A string in C is encoded by an array of character terminated by the null character.  
    **TRUE**    **FALSE**
- e. x86-64 instructions can range in length from 1 to 15 bytes.    **TRUE**    **FALSE**
- f. x86-64 is a 64-bit instruction set architecture.    **TRUE**    **FALSE**
- g. In x86-64 `%rsp` is the stack register and you can use an add instruction to add an immediate value to it.    **TRUE**    **FALSE**
- h. A benefit of a fully associative cache is that it tends to reduce the miss rate.    **FALSE**    **TRUE**
- i. In x86-64 `cmpq %rdx, %rax` would not alter the condition codes.    **FALSE**    **TRUE**
- j. In x86-64 `cmpq %rdx, %rax` would leave the value in `%rdx` unchanged.    **FALSE**    **TRUE**
- k. In x86-64 `cmpq %rdx, %rax` would leave the value in `%rax` unchanged.    **FALSE**    **TRUE**
- l. In x86-64, both `sar` and `shl` shift instructions have the same effect.    **FALSE**    **TRUE**

**1a:  $2^{(n-1)}$  negative numbers vs.  $2^{(n-1)}-1$  positive numbers.**

**1e: See Section 3.2.2, page 174 of the textbook.**

**1g: See Figure 3.31 of the textbook for an example.**

**1i, j, k: See Section 3.6.1 of the textbook.**

**1l: SAR and SHL are two names for the same instruction, see Section 3.5.3.**

2. In your own words, describe how parameters are passed to a procedure using the registers and the stack. (12 points)

Convention on the x86-64 ISA dictates that the first 6 integral (i.e., integers and pointers) arguments are passed through registers (rdi, rsi, rdx, rcx, r8, and r9), non-integral and additional arguments are passed through the stack, and the return value is passed through register rax.

If arguments are passed on the stack, values are pushed on the stack consecutively and in reverse order (argument 7 is at the top of the stack, i.e., at a lower address than argument 8), and all data sizes are rounded up to be multiple of 8 bytes.

3. Let  $x=0x87$ , give the hex for the following assuming a word size of 8 bits (6 Points)
- a.  $x \ll 3$

0x38

- b.  $x \gg 2$  (logical)

0x21

- c.  $x \gg 2$  (arithmetic)

0xe1

4. Illustrate the difference between the following two C functions by implementing them using Y86 assembly. (12 points)

```
long inc_a(long a) { return a++; }
```

```
long inc_b(long b) { return ++b; }
```

```
inc_a: irmovq $1, %r10    # constant 1
```

```
      addq %r10, %rdi    # a++
```

```
      rrmovq %rdi, %rax  # return the updated value
```

```
      ret
```

```
inc_b: rrmovq %rdi, %rax  # return the original value
```

```
      irmovq $1, %r10    # constant 1
```

```
      addq %r10, %rdi    # ++b
```

```
      ret
```

Since the value of the “++b” computation is not used, the corresponding instruction could be discarded.

5. Assume a non-pipelined processor with a combinational circuit that takes 300 ps and a register element that takes 20 ps to latch a new value. (12 points)

- a. What is the cycle time for the non-pipelined processor?

Cycle time =  $300 + 20 = 320\text{ps}$ , or

Throughput =  $1,000 [\text{ps/ns}] / 320 \text{ ps} = 3.13 \text{ Giga-instructions/s}$

- b. Suppose a computer architect proposes to break up the combinational circuit in 3 stages to build a pipelined version of the original design. Each stage takes 100 ps. A pipeline register follows each stage; each register taking 20 ps. What is the cycle time (after how much time can the next instruction start execution; not the total delay) of this 3-stage pipeline?

Cycle time =  $100 + 20 = 120\text{ps}$ , or

Throughput =  $1,000 [\text{ps/ns}] / 120 \text{ ps} = 8.33 \text{ Giga-instructions/s}$

- c. Suppose the architect cannot evenly break the original 300 ps delay among the 3 stages. Instead stage 1 takes 50 ps, stage 2 takes 150 ps, and stage 3 takes 100 ps. We are still using the same pipeline registers. What is the fastest cycle time that this design can achieve?

Cycle time =  $150 + 20 = 170\text{ps}$ , or

Throughput =  $1,000 [\text{ps/ns}] / 170 \text{ ps} = 5.88 \text{ Giga-instructions/s}$

The rate at which the pipeline can operate is limited by its slowest stage.

6. Consider the following x86-64 assembly code for a C loop. (12 points)

```
loop:
    movl    $0x0, %eax
    jmp     .L2
.L3:
    subl    $1, %edi
    addl    $1, %esi
    addl    $1, %eax
.L2:
    cmpl    %esi, %edi
    jg      .L3
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code.  
(Note: you may only use the symbolic variables `x`, `y`, and `result` in your expressions below — do not use register names.)

```
int loop(int x, int y)
{
    int result;

    for ( result = 0 ; x > y ; result++ ) {
        x = x - 1 ;
        y = y + 1 ;
    }

    return result;
}
```

7. Consider the following C functions and x86-64 assembly code. (6 points)

```
int fun1(int a, int b)
{
    if (a < b)
        return a;
    else
        return b;
}
```

```
int fun2(int a, int b)
{
    if (a > b)
        return a;
    else
        return b;
}
```

```
int fun3(int a, int b)
{
    if (a >= b)
        return a;
    else
        return b;
}
```

```
movl    %esi, %eax
cmpl    %esi, %edi
j1       .L3
.L2:
ret
.L3:
movl    %edi, %eax
jmp     .L2
```

Which of the C functions on the left represent the assembly code on the right?

8. The following problem concerns basic cache lookups. (16 points)

- The memory is byte addressable.
- Memory accesses are **1-byte words** (not 4-byte words).
- Memory addresses are 13-bits wide.
- The cache is 2-way set associative, with 4 byte block size and 16 total blocks.

In the following table, all numbers are given in hexadecimal. The contents of the cache are as follows:

2-way Set Associative Cache												
Index	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3
0	09	1	86	30	3F	10	00	0	99	04	03	48
1	45	1	60	4F	E0	23	38	1	00	BC	0B	37
2	EB	0	2F	81	FD	09	0B	0	8F	E2	05	BD
3	06	0	3D	94	9B	F7	32	1	12	08	7B	AD
4	C7	1	06	78	07	C5	05	1	40	67	C2	3B
5	71	1	0B	DE	18	4B	6E	0	B0	39	D3	F7
6	91	1	A0	B7	26	2D	F0	0	0C	71	40	10
7	46	0	B1	0A	32	0F	DE	1	12	C0	88	37

- a. The box below shows the format of a memory address. Indicate (by labeling the diagram) the fields that would be used to determine the following:

CO The block offset with the cache block  
 CI The cache index  
 CT The cache tag

12	11	10	9	8	7	6	5	4	3	2	1	0
CT	CT	CT	CT	CT	CT	CT	CT	CI	CI	CI	CO	CO

- b. For the given memory address, indicate the cache entry accessed and the cache byte value returned in hex. Indicate whether a cache miss occurs.

If there is a cache miss, enter “—” for “Cache Byte Returned.”

Memory address: **0xE34**

I. Memory address format (one bit per box):

12	11	10	9	8	7	6	5	4	3	2	1	0
0	1	1	1	0	0	0	1	1	0	1	0	0

II. Memory reference:

Parameter	Value
Block offset CO	0x 0
Cache Index CI	0x 5
Cache Tag CT	0x 71
Cache Hit? (Y/N)	Y
Cache Byte Returned	0x 0B