

1. For the following, circle **TRUE** or **FALSE** (24 points)

- a. Two's complement allows us to represent the same number of positive integers as negative integers for a given word length.     **TRUE**     **FALSE**
- b. If `%edx = 0xABABABAB` and `%eax = 0x4455AA22`, the x86-64 instruction  
    `movsbl %dh, %eax`  
    would change the value of `%rax` to `0xFFFFFFFFAB`     **TRUE**     **FALSE**
- c. The variables local to a function are allocated from the stack.     **TRUE**     **FALSE**
- d. A string in C is encoded by an array of character terminated by the null character.  
    **TRUE**     **FALSE**
- e. x86-64 instructions can range in length from 1 to 15 bytes.     **TRUE**     **FALSE**
- f. x86-64 is a 64-bit instruction set architecture.     **TRUE**     **FALSE**
- g. In x86-64 `%rsp` is the stack register and you can use an add instruction to add an immediate value to it.     **TRUE**     **FALSE**
- h. A benefit of a fully associative cache is that it tends to reduce the miss rate.     **TRUE**  
    **FALSE**
- i. In x86-64 `cmpq %rdx, %rax` would not alter the condition codes.     **TRUE**  
    **FALSE**
- j. In x86-64 `cmpq %rdx, %rax` would leave the value in `%rdx` unchanged.     **TRUE**  
    **FALSE**
- k. In x86-64 `cmpq %rdx, %rax` would leave the value in `%rax` unchanged.     **TRUE**  
    **FALSE**
- l. In x86-64, both `sal` and `shl` shift instructions have the same effect.     **TRUE**  
    **FALSE**

2. In your own words, describe how parameters are passed to a procedure using the registers and the stack. (12 points)

3. Let  $x=0x87$ , give the hex for the following assuming a word size of 8 bits (6 Points)

a.  $x \ll 3$

b.  $x \gg 2$  (logical)

c.  $x \gg 2$  (arithmetic)

4. Illustrate the difference between the following two C functions by implementing them using Y86 assembly. (12 points)

```
long inc_a(long a) { return a++; }  
long inc_b(long b) { return ++b; }
```

5. Assume a non-pipelined processor with a combinational circuit that takes 300 ps and a register element that takes 20 ps to latch a new value. (12 points)
- What is the cycle time for the non-pipelined processor?
  - Suppose a computer architect proposes to break up the combinational circuit in 3 stages to build a pipelined version of the original design. Each stage takes 100 ps. A pipeline register follows each stage; each register taking 20 ps. What is the cycle time (after how much time can the next instruction start execution; not the total delay) of this 3-stage pipeline?
  - Suppose the architect cannot evenly break the original 300 ps delay among the 3 stages. Instead stage 1 takes 50 ps, stage 2 takes 150 ps, and stage 3 takes 100 ps. We are still using the same pipeline registers. What is the fastest cycle time that this design can achieve?

6. Consider the following x86-64 assembly code for a C loop. (12 points)

```
loop:
    movl    $0x0, %eax
    jmp     .L2
.L3:
    subl    $1, %edi
    addl    $1, %esi
    addl    $1, %eax
.L2:
    cmpl    %esi, %edi
    jg      .L3
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code.  
(Note: you may only use the symbolic variables `x`, `y`, and `result` in your expressions below — do not use register names.)

```
int loop(int x, int y)
{
    int result;

    for ( _____; _____; result++ ) {
        _____;
        _____;
    }

    return result;
}
```

7. Consider the following C functions and x86-64 assembly code. (6 points)

<pre>int fun1(int a, int b) {     if (a &lt; b)         return a;     else         return b; }</pre>		
<pre>int fun2(int a, int b) {     if (a &gt; b)         return a;     else         return b; }</pre>	<pre>.L2     ret .L3</pre>	<pre>movl    %esi, %eax cmpl    %esi, %edi jl      .L3</pre>
<pre>int fun3(int a, int b) {     if (a &gt;= b)         return a;     else         return b; }</pre>		<pre>movl    %edi, %eax jmp     .L2</pre>

Which of the C functions on the left represent the assembly code on the right?

8. The following problem concerns basic cache lookups. (16 points)

- The memory is byte addressable.
- Memory accesses are **1-byte words** (not 4-byte words).
- Memory addresses are 13-bits wide.
- The cache is 2-way set associative, with 4 byte block size and 16 total blocks.

In the following table, all numbers are given in hexadecimal. The contents of the cache are as follows:

2-way Set Associative Cache												
Index	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3
0	09	1	86	30	3F	10	00	0	99	04	03	48
1	45	1	60	4F	E0	23	38	1	00	BC	0B	37
2	EB	0	2F	81	FD	09	0B	0	8F	E2	05	BD
3	06	0	3D	94	9B	F7	32	1	12	08	7B	AD
4	C7	1	06	78	07	C5	05	1	40	67	C2	3B
5	71	1	0B	DE	18	4B	6E	0	B0	39	D3	F7
6	91	1	A0	B7	26	2D	F0	0	0C	71	40	10
7	46	0	B1	0A	32	0F	DE	1	12	C0	88	37

- a. The box below shows the format of a memory address. Indicate (by labeling the diagram) the fields that would be used to determine the following:

CO    The block offset with the cache block  
CI    The cache index  
CT    The cache tag

12	11	10	9	8	7	6	5	4	3	2	1	0

- b. For the given memory address, indicate the cache entry accessed and the cache byte value returned in hex. Indicate whether a cache miss occurs.

If there is a cache miss, enter “–” for “Cache Byte Returned.”

Memory address: **0xE34**

I. Memory address format (one bit per box):

12	11	10	9	8	7	6	5	4	3	2	1	0

II. Memory reference:

Parameter	Value
Block offset CO	0x
Cache Index CI	0x
Cache Tag CT	0x
Cache Hit? (Y/N)	
Cache Byte Returned	0x